



AI COMMUNITY CONFERENCE

Arlington, VA, USA
September 11, 2026

Copilot Studio + Microsoft
Foundry:
When (and How) to Graduate Off
the Low-Code Path

#AICOEvents



Abhijeet Jadhav

- Director - Cloud Solutions, TrnDigital

6x Microsoft MVP.

PMP Certified Professional with 22+ years designing enterprise platforms.

LinkedIn:



Session Agenda

01 What Changed — The Platform Shift Since 2025

02 Platform Anatomy — Head-to-Head, August 2026

03 GitHub Copilot harness Deep Dive

04 The Revised Graduate Signals

05 Decision Framework + Live Demo

06 Architecture Patterns A / B / C / D

07 Governance, Cost & Migration

08 Q&A +Resources

What Changed — The Platform Shift Since 2025

Both platforms have moved dramatically. The old ceiling is gone; a new decision map is needed.

Copilot Studio — Leaps Forward

- GitHub Copilot harness GA (Aug 2026) — reasoning-first orchestration
- MCP servers as native tools — connect any external service
- Foundry IQ grounding — use Foundry knowledge from within Studio
- Persistent memory across sessions per user
- Reusable Skills — modular, shareable instruction sets
- Computer use tool — agents automate browsers & desktop apps
- Real-time voice agents — voice + digital from one agent
- Claude Sonnet 4.x / GPT-5 / Mistral selectable per agent
- Agent-to-agent (A2A) GA — delegate to specialist agents
- Agent Evaluations GA with CI/CD pipeline integration

Microsoft Foundry — Production-Ready

- Hosted Agents GA — managed infra, no server ops needed
- Foundry IQ — unified enterprise knowledge plane (OneLake, AI Search, MCP, Web IQ)
- Toolboxes in preview — curated tool bundles for agents
- Voice Live — real-time voice path in Foundry agents
- Publish Foundry agents to M365 Copilot and Teams
- Agent Optimizer — score, tune, and version agents via CLI
- OpenTelemetry-compatible tracing across model calls, tool invocations, and sub-agent hops
- Foundry model catalog — discover, purchase, and integrate models
- GPT-5 RFT (gated GA), Grok 4.3, DeepSeek V4, MAI models
- Private Agentic Retrieval — retrieval inside private network boundary

What Stayed True

- Studio = low-code speed; Foundry = full-code control
- Studio Credits vs Azure consumption cost model unchanged
- Foundry still requires a developer team to build & operate
- Hybrid (Studio UX + Foundry brain) remains the power pattern
- Compliance-critical immutable audit trails still need Foundry
- Version-pinned model deployment still a Foundry differentiator
- Complex deterministic orchestration still stronger in Foundry
- Migration cost is real — don't over-engineer for Studio limits

Platform Anatomy — Head-to-Head, August 2026

✓ Aug 2026

	Copilot Studio	Microsoft Foundry (Azure AI Foundry)
Audience	Makers, Power Users, Pro Devs (harness)	Developers & ML Engineers
Model Control	GPT-4.1, GPT-5 Chat, Claude 4.x/Sonnet 5, Mistral, GPT-5.5 Chat (selectable per agent; external models may process outside tenant)	Any model, any version, version-pinned; 10,000+ models in Foundry model catalog; GPT-5 RFT, Grok 4.3, DeepSeek V4 available
Orchestration	Standard harness (topic flows) OR GitHub Copilot harness (GA Aug 2026) — reasoning, breaks goals into steps, self-heals, MCP, skills	Foundry Agent Service (Hosted Agents GA), Semantic Kernel, LangChain, LangGraph, Agent Optimizer CLI for iterative tuning
Grounding	SharePoint, Dataverse, Azure AI Search, Foundry IQ (connect Foundry KB from Studio), Work IQ (email, calendar, Teams, files)	Foundry IQ — unified knowledge plane; OneLake, AI Search, MCP sources, Web IQ, Private Agentic Retrieval (in-VNet)
Memory	Persistent per-user memory across sessions (GitHub Copilot harness agents, June 2026)	Procedural, user & session memory in Foundry Agent Service (public preview)
Observability	Power Platform analytics + App Insights (env-level export to App Insights, July 2026) Agent Evaluations GA + CI/CD integration	Full OpenTelemetry — every model call, tool invoke, sub-agent hop; trace-based evals; Agent Optimizer links evals back to traces
Auth & Identity	Entra ID built-in; Entra Agent ID per agent (mandatory July 2026); scoped DLP + Conditional Access	Full Entra + custom auth flows; scoped per Foundry project & deployment
Cost Model	Copilot Credits (complexity-dependent; M365-licensed users = zero credits for classic)	Azure consumption (token-based, compute, AI Search); no seat licence fee
Time to Deploy	Hours to days	Days to weeks (+ ongoing DevOps)

GitHub Copilot harness — The Game Changer

GA August 2026 · Reasoning-first · Built for complex, multi-step business processes

GitHub Copilot harness (GA)

Skills

Reusable modular instruction sets

Memory

Persistent per-user context across sessions

MCP Servers

Any external service as a native tool (GA)

Connected Agents

Delegate to specialist agents (A2A GA)

Foundry IQ

Ground in Foundry knowledge bases

Work IQ

Email, calendar, Teams files & people data

Computer Use

Automate browser & desktop apps

Workflows + Files

Multistep automation; create Word/Excel/PDF



Harness agents can take a goal, break it into steps, call tools across connectors / knowledge / MCP / connected agents, and retries/finds alternative paths automatically — not script-bound

The Revised Graduate Signals

Recalibrated for August 2026 — Studio has closed several gaps. These 5 signals still require Foundry.

01

Version-Pinned Custom Model Deployment

You need a specific model version that never changes under you across platform updates, OR you need Reinforcement Fine-Tuning (GPT-5 RFT), full custom fine-tuning, or a model not in Studio's catalog. Studio selects models but does not pin versions or run your own fine-tuned weights.

💡 Studio model selection ≠ version pinning. Fine-tuning still requires Foundry.

02

Sovereign RAG or Extreme-Scale Knowledge

Your knowledge base exceeds ~500 documents with complex retrieval needs AND you cannot use Foundry IQ grounding from within Studio. Studio now supports Foundry IQ as a grounding source — but if you need Private Agentic Retrieval (in-VNet), hybrid vector+keyword with custom chunking, or 50K+ docs from proprietary stores, Foundry's retrieval stack is still required.

💡 Foundry IQ grounding in Studio partially closes this gap. Reassess before migrating.

03

Immutable Compliance Audit Trail

Legal/security requires immutable, tamper-evident logs of every prompt, response, and data retrieval — not just usage telemetry. Studio's App Insights export (July 2026) improves observability, but logs are still split across Power Platform and Azure. Full OpenTelemetry pipeline end-to-end, linking evals to traces, is a Foundry-only capability.

💡 Studio now exports env-level telemetry, but end-to-end immutable trace is still Foundry.

04

Deterministic Pro-Code Orchestration

Your agent needs code-level control over orchestration loops — branching logic that cannot be expressed as instructions, custom retry policies, complex state machines, or integration with existing CI/CD and DevOps pipelines via the Agent Optimizer CLI. Studio's harness is powerful but instruction-driven; Foundry gives you full SDK control.

💡 MCP + A2A in Studio closes many orchestration gaps. This signal now requires a tighter test.

05

Cost at Scale with Complex Reasoning

Copilot Credit costs exceed Azure consumption pricing — typically at ~50K+ monthly interactions using reasoning models. A GPT-5.5 Reasoning or Claude Opus call in Studio can cost 50× a classic answer. Foundry's token-based billing is significantly cheaper for reasoning-heavy, high-volume workloads.

💡 Use the Studio Agent Usage Estimator (April 2026) before migrating on cost alone.

The Decision Framework

Answer each question about your scenario — not about what platforms can do

How to use: tick the box that matches YOUR scenario. Count your Foundry ticks.

Signal / Question	Migration Pressure	My need is Low or None	My need is Moderate	Migration pressure key: My need is High → signals Foundry
Model version & fine-tuning Do I need to pin a specific model version that never changes, or run custom fine-tuned weights (e.g. GPT-5 RFT)?	HIGH	No — model selection per prompt is enough → Studio or Harness	Somewhat — I want stability but can tolerate platform-managed updates → Studio + Harness	Yes — I must pin a version or deploy fine-tuned weights → Microsoft Foundry
Knowledge scale & sovereignty Does my knowledge base exceed ~500 docs, require in-VNet retrieval, or use sources outside SharePoint/Dataverse?	HIGH	No — SharePoint & Dataverse grounding is enough → Studio	Moderate — I need Foundry IQ grounding from within Studio → Studio + Harness	Yes — 50K+ docs, custom chunking, or Private Agentic Retrieval needed → Microsoft Foundry
Compliance & audit trail Does legal/security require a tamper-evident, end-to-end trace of every prompt, response, and retrieval event?	HIGH	No — Power Platform analytics is sufficient → Studio	Partial — App Insights export covers our audit needs → Studio + Harness	Yes — we need full immutable OTel pipeline linking evals to traces → Microsoft Foundry
Orchestration complexity Does my agent need code-level orchestration logic, custom retry/state machines, or CI/CD-gated versioning?	HIGH	No — topic flows are sufficient → Studio	Moderate — harness orchestration + MCP covers my needs → Studio + Harness	Yes — I need SDK-level control, Agent Optimizer CI/CD, or deterministic loops → Microsoft Foundry
Dev team availability Do I have developers who can write Python or C# and operate Azure infrastructure?	MED	No — makers and power users only → Studio	Low-code — some tech capability but not full-stack devs → Studio + Harness	Yes — we have Python or C# developers with Azure experience → Foundry is viable
Monthly interaction volume & cost Do my Copilot Credit costs exceed the Azure consumption equivalent at my interaction volume and complexity?	MED	< 20K classic interactions / month → Studio is cost-optimal	20K–100K reasoning interactions — near crossover point → Model the cost	100K+ reasoning interactions — Foundry's token model is cheaper → Microsoft Foundry

0–1 Foundry ticks → Stay in Copilot Studio	2 Foundry ticks → Studio + GitHub Copilot harness	3 Foundry ticks → Hybrid (Pattern C)	4+ Foundry ticks → Microsoft Foundry (Pattern D)
--	---	--------------------------------------	--

Architecture Patterns — Four in 2026

Four patterns from no-code to full pro-code — pick the one that matches your signal score

A — Copilot Studio Only

Score 0

Low complexity

- Copilot Studio standard harness
- SharePoint / Dataverse grounding
- Entra ID auth (built-in)
- Purview DLP + Entra Agent ID

Best for: HR bots, FAQ agents, simple approvals

B — Studio + Copilot Harness

Score 1–2

Medium complexity

- GitHub Copilot harness (GA)
- MCP servers as tools
- Skills + persistent memory
- Foundry IQ or Work IQ grounding
- A2A agent delegation

Best for: Complex business process automation, document workflows, multi-step reasoning agents

C — Studio Front-end, Foundry Brain

Score 2–3

Hybrid

- Studio (UX + channels + Entra Agent ID)
- Foundry brain — Hosted Agents
- Foundry IQ retrieval plane
- Shared Entra identity + DLP
- Full OTel tracing from Foundry

Best for: Customer-facing agents, compliance teams with M365 channel requirements

D — Full Microsoft Foundry

Score 3+

High complexity

- Microsoft Foundry Agent Service (Hosted GA)
- Semantic Kernel or LangChain SDK
- Agent Optimizer + CI/CD
- Full OTel + immutable audit
- Publish to Teams & M365 via Activity protocol

Best for: Compliance-critical, ML-intensive, sovereign RAG, version-pinned models

Governance & Cost Trade-offs

Cost scales with complexity per interaction — not just volume. Studio's observability has improved significantly.

Governance Capabilities

Capability	Studio	Studio + Harness	Foundry
DLP Policies	✓	✓	✓
Sensitivity Labels	✓	✓	✓
Entra Agent ID (per agent)	✓ Jul 2026	✓	✓
Prompt Audit Log	Partial*	Partial*	✓ OTel
Model Version Lock	✗	✗	✓
Agent Evaluations (CI/CD)	✓ GA Mar 2026	✓	✓
Safety Evaluations	✗	✗	✓
SOC 2 / ISO 42001	Partial	Partial	✓
Immutable End-to-End Trace	✗	✗	✓ OTel

* Prompt Audit Log now available via Power Platform admin center App Insights export (July 2026) — setup still required

Real-World Cost Scenarios

Cost scales with complexity per interaction — not just volume

Scenario	Studio est.	Foundry est.
HR FAQ bot 8K interactions/mo M365-licensed, classic	\$0–\$80 (M365 licensed = zero credits)	\$180–\$280 (Azure OpenAI + infra overhead)
Business process agent 40K/mo · external users Harness + reasoning	\$600–\$900 (avg 4–8 credits; reasoning model)	\$380–\$520 (near crossover point)
Compliance doc analysis 75K/mo · complex RAG Graph + Foundry IQ	\$2,800–\$4,500 (12–22 credits avg per interaction)	\$900–\$1,400 (token-based consumption)

Key: Studio costs hinge on complexity per interaction — not just volume. Use the Studio Agent Usage Estimator before migrating on cost alone.

Migration Walkthrough — Classic Studio → Harness or Hybrid

Two migration paths: upgrade to Harness (Pattern B) or go full Hybrid/Foundry (Pattern C/D)

PATH 1 — Upgrade to GitHub Copilot harness (weeks)

Step 1

Day 1

Enable Harness

Create a new agent in Copilot Studio and select the GitHub Copilot harness. Re-implement behavior using instructions — harnesses cannot be transferred between each other.

Step 2

Day 1–3

Migrate Topics → Instructions

Convert topic flows to natural-language instructions. Harness understands goals, not scripts.

Step 3

Day 3–7

Add MCP + Skills

Wire MCP servers as tools. Create/upload reusable Skills for common sub-tasks.

Step 4

Day 5–8

Enable Memory + Foundry IQ

Turn on persistent memory. Connect Foundry IQ knowledge base if external grounding needed.

Step 5

Day 7–14

Evaluate + Rollout

Run agent evaluations via CI/CD. Compare quality metrics vs classic harness. Cut over.

PATH 2 — Full Hybrid / Foundry Migration (2–4 weeks)

Step 1

Day 1

Export Agent Definition

Extract topics, knowledge sources, and connectors from Studio as solution ZIP.

Step 2

Day 1–2

Provision Foundry Hosted Agent

Create Foundry project, configure Azure OpenAI deployment, pin model version. Enable Hosted Agents.

Step 3

Day 3–7

Rebuild Orchestration

Port topic flow logic to Semantic Kernel agent or Foundry Hosted Agent with tool calling.

Step 4

Day 5–9

Connect Channels + Observability

Publish Foundry agent to Teams via Activity protocol and Microsoft 365 publish APIs. Wire App Insights, OTel, Purview DLP on Day 1.

Step 5

Day 10–14

Parallel Run, Evaluate & Cutover

Run both agents simultaneously. Compare quality via Foundry Evaluations. Rollback trigger: p95 latency >30% above Studio baseline for 48h.

Key Takeaways — August 2026 Edition

The platform has moved. Your decision framework should too.

1

Studio is no longer just a starter tier

With the GitHub Copilot harness GA, MCP, persistent memory, and Foundry IQ grounding, Studio handles use cases that previously required Foundry. Re-evaluate your agents before migrating.

2

The harness changes the graduate threshold

Signals 02 (RAG) and 04 (orchestration) now have a 'Harness-first' answer before you go to Foundry. Test the harness against your scenario before committing to migration.

3

Hybrid is still the power pattern — but it's cleaner

Foundry IQ in Studio means you can ground in Foundry knowledge without a full migration. The hybrid boundary is now Studio Harness (UX + reasoning) + Foundry (sovereign RAG + model control).

4

Foundry's differentiators have sharpened

Version-pinned fine-tuned models, full OTel immutable audit trail, Agent Optimizer CLI, and Private Agentic Retrieval are the clear Foundry-only capabilities in 2026. Lead with these in your go/no-go.

5

Governance is more equal — but immutability is still Foundry

Studio now has Entra Agent IDs, App Insights export, and GA evaluations. The remaining gap is the end-to-end immutable OTel trace that links every eval back to the exact model call that produced it.

Thank You

Questions & Resources

→ **Copilot Studio What's New**

learn.microsoft.com/en-us/microsoft-copilot-studio/whats-new

→ **Copilot Studio Harnesses Overview**

learn.microsoft.com/en-us/microsoft-copilot-studio/harnesses-overview

→ **Microsoft Foundry Blog**

devblogs.microsoft.com/foundry

→ **Foundry Agent Service Docs**

learn.microsoft.com/en-us/azure/ai-services/agents/overview

→ **Copilot Studio Security Top 10**

aka.ms/CopilotStudioSecurity

→ **Copilot Agent Kit (Power CAT)**

github.com/microsoft/Power-CAT-Copilot-Studio-Kit

→ **MCP in Copilot Studio**

learn.microsoft.com/en-us/microsoft-copilot-studio/mcp-add-components-to-agent

→ **Agent Usage Estimator**

learn.microsoft.com/en-us/microsoft-copilot-studio/agent-usage-estimator

Upcoming Local Events/Conferences

- **AI-ML User Group**
 - Monthly In-Person Meetings
 - Reston, Virginia, USA
 - <https://www.meetup.com/artificialintelligenceandmachinelearning/>



Speaker Evaluations

Please complete an evaluation for each session you are attending. This will help us in providing our speakers with valuable feedback. You will also increase your chances of winning a raffle prizes



Thank you for attending

We hope to see you at one of our upcoming events



AI Community Conference DC 2026
#AICommunityConference